

Computer Systems and Networks

Prof. Forsyth

Homework 2 - Assembly Review

Due: September 12th 2023

1 Problem 1: Getting Started with the LC-2222

In this homework, you will be using the LC-2222 ISA to complete a Tower of Hanoi move-counting function. Before you begin, you should familiarize yourself with the available instructions, the register conventions and the calling convention of LC-2222. Details can be found in the section, Appendix A: LC-2222 Instruction Set Architecture, at the end of this document.

The `assembly` folder contains several tools for you to use:

- `assembler.py`: a basic assembler that can take your assembly code and convert it into binary instructions for the LC-2222.
- `lc2222.py`: the ISA definition file for the assembler, which tells `assembler.py` the instructions supported by the LC-2222 and their formats.
- `lc2222-sim.py`: A simulator of the LC-2222 machine. The simulator reads binary instructions and emulates the LC-2222 machine, letting you single-step through instructions and check their results.

To learn how to run these tools, see the `README.md` file in the `assembly` directory.

Before you begin work on the second problem of the homework, try writing a simple program for the LC-2222 architecture. This should help you familiarize yourself with the available instructions.

We have provided a template, `mod.s`, for you to use for this purpose. Try writing a program that performs the `mod` operation on the two provided arguments. A correct implementation will result in a value of 2.

You can use the following C code snippet as a guide to implement this function:

```
int mod(int a, int b) {
    int x = a;
    while (x >= b) {
        x = x - b;
    }
    return x;
}
```

There is no turn-in for this portion of the assignment, but it is **recommended** that you attempt it in order to familiarize yourself with the ISA.

2 Problem 2: Tower of Hanoi

For this problem, you will be implementing the missing portions of the program that calculates the minimum number of moves to solve the Tower of Hanoi problem for n disks.

Tower of Hanoi involves three vertical rods and a set of varying sized disks, which can slide onto any rod. The disks are initially stacked on one of the rods in ascending order of size, with the largest disk on the bottom and the smallest on top, thus making a conical shape. The objective of this puzzle is to migrate the tower of disks completely to another rod, under the rule that only individual disks may be moved at once, and no disks may be placed on smaller disks.

You will be finishing a **recursive** implementation of the Tower of Hanoi minimal moves calculator program that follows the LC-2222 calling convention. Recursive functions always obtain a return address through the function call and return to the callee using the return address.

You must use the stack pointer (\$sp) and frame pointer (\$fp) registers as described in the textbook and lecture slides.

Here is the C code for the Tower of Hanoi minimal moves calculator you have been provided:

```
int minimumHanoi(int n) {
    if (n == 1)
        return 1;
    else
        return (2 * minimumHanoi(n - 1)) + 1;
}
```

Note that this C code is just to help your understanding and does not need to be exactly followed. However, your assembly code implementation should meet all of the given conditions in the description.

Open `hanoi.s` file in the assembly directory. This file contains an implementation of the Tower of Hanoi minimal moves calculator program that is missing significant portions of the calling convention. Near the bottom of the `hanoi.s` we have provided multiple numbers that you can use to test your homework. They are located at labels `testNumDisks1`, `testNumDisks2`, `testNumDisks3`. Be sure to use these provided integers by loading them from the labels into registers. None of the numbers provided and tested will be lower than 1.

Complete the program by implementing the various missing portions of the LC-2222 calling convention. Each location where you need to implement a portion of the calling convention is marked with a `TODO` label as well as a short hint describing the portion of the calling convention you should be implementing.

Please note that we will be testing your implementation for multiple different instances, so please do not attempt to hardcode your solutions.

3 Problem 3: Short Answer

Please answer the following question in the file named `answers.txt`:

1. The LC-2222 instruction set contains an instruction called `jalr` that is used to jump to a location while saving a return address. However, this functionality could be emulated using a combination of other instructions available in the ISA. Describe a sequence of other instructions in the LC-2222 ISA that you may use to accomplish the functionality of `jalr`.

For the purpose of this question, you may assume the target address is represented with the label `<target>` which can be accessed using the 20 bits reserved for an offset or immediate value in the LC-2222 ISA.

4 Deliverables

- `hanoi.s`: your assembly code from Section 2
- `answers.txt`: your answer to the problem from Section 3

Submit these files to **Gradescope** before the assignment deadline.

The TAs should be able to type `python assembler.py -i lc2222 --sym hanoi.s` and then `python lc2222-sim.py hanoi.bin` to run your code. If you cannot do this with your submission, then you have done something wrong.

5 Appendix A: LC-2222 Instruction Set Architecture

The LC-2222 is a simple, yet capable computer architecture. The LC-2222 combines attributes of both ARM and the LC-2200 ISA defined in the Ramachandran & Leahy textbook for CS 2200.

The LC-2222 is a **word-addressable, 32-bit** computer. **All addresses refer to words**, i.e. the first word (four bytes) in memory occupies address 0x0, the second word, 0x1, etc.

All memory addresses are truncated to 16 bits on access, discarding the 16 most significant bits if the address was stored in a 32-bit register. This provides roughly 64 KB of addressable memory.

5.1 Registers

The LC-2222 has 16 general-purpose registers. While there are no hardware-enforced restraints on the uses of these registers, your code is expected to follow the conventions outlined below.

Table 1: Registers and their Uses

| Register Number | Name | Use | Callee Save? |
|-----------------|--------|---------------------------|--------------|
| 0 | \$zero | Always Zero | NA |
| 1 | \$at | Assembler/Target Address | NA |
| 2 | \$v0 | Return Value | No |
| 3 | \$a0 | Argument 1 | No |
| 4 | \$a1 | Argument 2 | No |
| 5 | \$a2 | Argument 3 | No |
| 6 | \$t0 | Temporary Variable | No |
| 7 | \$t1 | Temporary Variable | No |
| 8 | \$t2 | Temporary Variable | No |
| 9 | \$s0 | Saved Register | Yes |
| 10 | \$s1 | Saved Register | Yes |
| 11 | \$s2 | Saved Register | Yes |
| 12 | \$k0 | Reserved for OS and Traps | NA |
| 13 | \$sp | Stack Pointer | No |
| 14 | \$fp | Frame Pointer | Yes |
| 15 | \$ra | Return Address | No |

1. **Register 0** is always read as zero. Any values written to it are discarded. **Note:** for the purposes of this project, you must implement the zero register. Regardless of what is written to this register, it should always output zero.
2. **Register 1** is used to hold the target address of a jump. It may also be used by pseudo-instructions generated by the assembler.
3. **Register 2** is where you should store any returned value from a subroutine call.
4. **Registers 3 - 5** are used to store function/subroutine arguments. **Note:** registers 2 through 8 should be placed on the stack if the caller wants to retain those values. These registers are fair game for the callee (subroutine) to trash.
5. **Registers 6 - 8** are designated for temporary variables. The caller must save these registers if they want these values to be retained.
6. **Registers 9 - 11** are saved registers. The caller may assume that these registers are never tampered with by the subroutine. If the subroutine needs these registers, then it should place them on the stack and restore them before they jump back to the caller.
7. **Register 12** is reserved for handling interrupts. While it should be implemented, it otherwise will not have any special use on this assignment.

8. **Register 13** is the everchanging top of the stack; it keeps track of the top of the activation record for a subroutine.
9. **Register 14** is the anchor point of the activation frame. It is used to point to the first address on the activation record for the currently executing process.
10. **Register 15** is used to store the address a subroutine should return to when it is finished executing.

5.2 Instruction Overview

The LC-2222 supports a variety of instruction forms, only a few of which we will use for this project. The instructions we will implement in this project are summarized below.

Table 2: LC-2222 Instruction Set

| | 31 | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23 | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 | |
|------|------|----|----|----|----|----|----|----|-----|----|----|----|----|----|----|----|----|----|----|----|----|----|---|---|---|---|---|---|---|---|---|---|------------|
| ADD | 0000 | | | | | | | | DR | | | | | | | | | | | | | | | | | | | | | | | | SR2 |
| NAND | 0001 | | | | | | | | DR | | | | | | | | | | | | | | | | | | | | | | | | SR2 |
| ADDI | 0010 | | | | | | | | DR | | | | | | | | | | | | | | | | | | | | | | | | immval20 |
| LW | 0011 | | | | | | | | DR | | | | | | | | | | | | | | | | | | | | | | | | offset20 |
| SW | 0100 | | | | | | | | SR | | | | | | | | | | | | | | | | | | | | | | | | offset20 |
| BEQ | 0101 | | | | | | | | SR1 | | | | | | | | | | | | | | | | | | | | | | | | offset20 |
| JALR | 0110 | | | | | | | | AT | | | | | | | | | | | | | | | | | | | | | | | | RA |
| HALT | 0111 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | unused |
| BLT | 1000 | | | | | | | | SR1 | | | | | | | | | | | | | | | | | | | | | | | | SR2 |
| LEA | 1001 | | | | | | | | DR | | | | | | | | | | | | | | | | | | | | | | | | PCoffset20 |
| BGT | 1010 | | | | | | | | SR1 | | | | | | | | | | | | | | | | | | | | | | | | offset20 |
| OR | 1011 | | | | | | | | DR | | | | | | | | | | | | | | | | | | | | | | 0 | | SR2 |
| XOR | 1011 | | | | | | | | DR | | | | | | | | | | | | | | | | | | | | | | 1 | | SR2 |

5.2.1 Conditional Branching

Branching in the LC-2222 ISA is a bit different than usual. We have a set of branching instructions including BEQ, BLT, and BGT, which offer the ability to branch upon a certain condition being met. These instructions use comparison operators, comparing the values of two source registers. If the comparisons are true (for example, with the BGT instruction, if $SR1 > SR2$), then we will branch to the target destination of $incrementedPC + offset20$.

5.3 Detailed Instruction Reference

5.3.1 ADD

Assembler Syntax

ADD DR, SR1, SR2

Encoding

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|------|----|----|----|----|----|----|----|-----|----|----|----|--------|----|----|----|----|----|----|----|----|----|---|---|---|---|---|---|-----|---|---|---|
| 31 | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23 | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| 0000 | | | | DR | | | | SR1 | | | | unused | | | | | | | | | | | | | | | | SR2 | | | |

Operation

DR = SR1 + SR2;

Description

The ADD instruction obtains the first source operand from the SR1 register. The second source operand is obtained from the SR2 register. The second operand is added to the first source operand, and the result is stored in DR.

5.3.2 NAND

Assembler Syntax

NAND DR, SR1, SR2

Encoding

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|------|----|----|----|----|----|----|----|-----|----|----|----|--------|----|----|----|----|----|----|----|----|----|---|---|---|---|---|---|-----|---|---|---|
| 31 | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23 | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| 0001 | | | | DR | | | | SR1 | | | | unused | | | | | | | | | | | | | | | | SR2 | | | |

Operation

DR = ~(SR1 & SR2);

Description

The NAND instruction performs a logical NAND (AND NOT) on the source operands obtained from SR1 and SR2. The result is stored in DR.

HINT: A logical NOT can be achieved by performing a NAND with both source operands the same.

For instance,

NAND DR, SR1, SR1

...achieves the following logical operation: $DR \leftarrow \overline{SR1}$.

5.3.3 ADDI

Assembler Syntax

ADDI DR, SR1, immval20

Encoding

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|------|----|----|----|----|----|----|----|-----|----|----|----|----------|----|----|----|----|----|----|----|----|----|---|---|---|---|---|---|---|---|---|---|
| 31 | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23 | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| 0010 | | | | DR | | | | SR1 | | | | immval20 | | | | | | | | | | | | | | | | | | | |

Operation

DR = SR1 + SEXT(immval20);

Description

The ADDI instruction obtains the first source operand from the SR1 register. The second source operand is obtained by sign-extending the immval20 field to 32 bits. The resulting operand is added to the first source operand, and the result is stored in DR.

5.3.4 LW

Assembler Syntax

LW DR, offset20(BaseR)

Encoding

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|------|----|----|----|----|----|----|----|-------|----|----|----|----------|----|----|----|----|----|----|----|----|----|---|---|---|---|---|---|---|---|---|---|
| 31 | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23 | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| 0011 | | | | DR | | | | BaseR | | | | offset20 | | | | | | | | | | | | | | | | | | | |

Operation

DR = MEM[BaseR + SEXT(offset20)];

Description

An address is computed by sign-extending bits [19:0] to 32 bits and then adding this result to the contents of the register specified by bits [23:20]. The 32-bit word at this address is loaded into DR.

5.3.5 SW

Assembler Syntax

SW SR, offset20(BaseR)

Encoding

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|------|----|----|----|----|----|----|----|-------|----|----|----|----------|----|----|----|----|----|----|----|----|----|---|---|---|---|---|---|---|---|---|---|
| 31 | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23 | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| 0100 | | | | SR | | | | BaseR | | | | offset20 | | | | | | | | | | | | | | | | | | | |

Operation

MEM[BaseR + SEXT(offset20)] = SR;

Description

An address is computed by sign-extending bits [19:0] to 32 bits and then adding this result to the contents of the register specified by bits [23:20]. The 32-bit word obtained from register SR is then stored at this address.

5.3.6 BEQ

Assembler Syntax

BEQ SR1, SR2, offset20

Encoding

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|------|----|----|----|-----|----|----|----|-----|----|----|----|----------|----|----|----|----|----|----|----|----|----|---|---|---|---|---|---|---|---|---|---|
| 31 | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23 | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| 0101 | | | | SR1 | | | | SR2 | | | | offset20 | | | | | | | | | | | | | | | | | | | |

Operation

```
if (SR1 == SR2) {
    PC = incrementedPC + offset20
}
```

Description

A branch is taken if SR1 is equal to SR2. If this is the case, the PC will be set to the sum of the incremented PC (since we have already undergone fetch) and the sign-extended offset[19:0].

5.3.7 JALR

Assembler Syntax

JALR AT, RA

Encoding

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|------|----|----|----|----|----|----|----|----|----|----|----|--------|----|----|----|----|----|----|----|----|----|---|---|---|---|---|---|---|---|---|---|
| 31 | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23 | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| 0110 | | | | AT | | | | RA | | | | unused | | | | | | | | | | | | | | | | | | | |

Operation

RA = PC;

PC = AT;

Description

First, the incremented PC (address of the instruction + 1) is stored into register RA. Next, the PC is loaded with the value of register AT, and the computer resumes execution at the new PC.

5.3.8 HALT

Assembler Syntax

HALT

Encoding

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|------|----|----|----|--------|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|---|---|---|---|---|---|---|---|---|---|
| 31 | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23 | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| 0111 | | | | unused | | | | | | | | | | | | | | | | | | | | | | | | | | | |

Description

The machine is brought to a halt and executes no further instructions.

5.3.9 BLT

Assembler Syntax

BLT SR1, SR2, offset20

Encoding

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|------|----|-----|----|-----|----|----------|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|---|---|---|---|---|---|---|---|---|---|
| 31 | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23 | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| 1000 | | SR1 | | SR2 | | offset20 | | | | | | | | | | | | | | | | | | | | | | | | | |

Operation

```
if (SR1 < SR2) {
    PC = incrementedPC + offset20
}
```

Description

A branch is taken if SR1 is less than SR2. If this is the case, the PC will be set to the sum of the incremented PC (since we have already undergone fetch) and the sign-extended offset[19:0].

5.3.10 LEA**Assembler Syntax**

LEA DR, label

Encoding

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|------|----|----|----|----|----|----|----|--------|----|----|----|------------|----|----|----|----|----|----|----|----|----|---|---|---|---|---|---|---|---|---|---|
| 31 | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23 | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| 1001 | | | | DR | | | | unused | | | | PCoffset20 | | | | | | | | | | | | | | | | | | | |

Operation

DR = PC + SEXT(PCoffset20);

Description

An address is computed by sign-extending bits [19:0] to 32 bits and adding this result to the incremented PC (address of instruction + 1). It then stores the computed address into register DR.

5.3.11 BGT**Assembler Syntax**

BGT SR1, SR2, offset20

Encoding

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|------|----|----|----|-----|----|----|----|-----|----|----|----|----------|----|----|----|----|----|----|----|----|----|---|---|---|---|---|---|---|---|---|---|
| 31 | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23 | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| 1010 | | | | SR1 | | | | SR2 | | | | offset20 | | | | | | | | | | | | | | | | | | | |

Operation

```
if (SR1 > SR2) {
    PC = incrementedPC + offset20
}
```

Description

A branch is taken if SR1 is greater than SR2. If this is the case, the PC will be set to the sum of the incremented PC (since we have already undergone fetch) and the sign-extended offset[19:0].

5.3.12 OR**Assembler Syntax**

OR DR, SR1, SR2

Encoding

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|------|----|----|----|----|----|----|----|-----|----|----|----|--------|----|----|----|----|----|----|----|----|----|---|---|---|---|-----|---|---|---|---|---|
| 31 | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23 | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| 1011 | | | | DR | | | | SR1 | | | | unused | | | | | | | | | | | | 0 | | SR2 | | | | | |

Operation

DR = SR1 | SR2;

Description

The OR instruction obtains the first source operand from the SR1 register. The second source operand is obtained from the SR2 register. Perform the OR operation on the two operands, and the result is stored in DR.

5.3.13 XOR**Assembler Syntax**

XOR DR, SR1, SR2

Encoding

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|------|----|----|----|----|----|----|-----|----|----|--------|----|----|----|----|----|----|----|----|----|----|----|---|---|---|-----|---|---|---|---|---|---|
| 31 | 30 | 29 | 28 | 27 | 26 | 25 | 24 | 23 | 22 | 21 | 20 | 19 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| 1011 | | | | DR | | | SR1 | | | unused | | | | | | | | | | | | | | 1 | SR2 | | | | | | |

Operation

DR = SR1 (XOR) SR2;

Description

The XOR instruction obtains the first source operand from the SR1 register. The second source operand is obtained from the SR2 register. Perform the XOR operation on the two operands, and the result is stored in DR.